

000
001
002
003
004
005
006
007
008
009
010
011
012
013
014
015
016
017
018
019
020
021
022
023
024
025
026
027
028
029
030
031
032
033
034
035
036
037
038
039
040
041
042
043
044
045
046
047
048
049
050
051
052
053054
055
056
057
058
059
060
061
062
063
064
065
066
067
068
069
070
071
072
073
074
075
076
077
078
079
080
081
082
083
084
085
086
087
088
089
090
091
092
093
094
095
096
097
098
099
100
101
102
103
104
105
106
107

CSCI 1430 Final Project Report: License Plate Detector

Anonymous CVPR submission

Paper ID ****

1. Introduction

Our project aims to build a computer vision program that takes in images of cars with different backgrounds and orientations, detects (and outlines) the license plate, and reads its number if visible. Our main tasks were building a license plate detector using edge detection and feature extraction, then training a neural network model to read the plate input.

This task is challenging because of the nature of the kinds of images that can be provided. License plates have a variety of formats across different areas and countries, and coupling this with the fact that cars can further be in different orientations or visibility means our approach needs to accommodate a lot of variation in input. Our data comes from datasets with images of cars from the US, Europe, and Brazil, with annotated license plates.

Our project could have an important application in traffic systems. For instance, many toll systems could be improved by simply using CV to identify license plates rather than stopping each car. Furthermore, this has applications to law enforcement and vehicle management by the government.

2. Related Work

We used some online references and OpenCV tutorials for our work on edge detection and character segmentation. Our datasets were from prior studies on license plates. The neural network was developed without reference to external works.

3. Method

Our approach and code involved several components: license plate detection, license plate digit segmentation, and a forward pass through a neural network.

3.1. Datasets

Various forms of data were needed to gather the necessary information on license plates. There were three main datasets. The first dataset was pictures of cars with an arbitrary background, which was used for testing the bounding boxes our

detector returns. The second was pictures of cropped license plates that were used for testing our segmentation algorithm that divided (i.e. segmented) a license plate picture into single characters. Once the segmentation script was complete, we used it to save many cropped images of single characters from license plates. These were used to train our neural network for character classification and can be considered our third dataset.

3.2. Plate Detector

Our approach utilizes several image filters, contours from OpenCV, and our segmentation algorithm to eliminate false positives. First, given an image of a car with an arbitrary background and visible license plate, we convert the image to gray scale, apply a Gaussian blur, and use a vertical Sobel filter to highlight vertical edges. We then use OpenCV's findContour method to find possible bounding boxes where all pixels on the border of a bounding box have nearly the same intensity. This gives us many possible bounding boxes, the cropped license plate being one of them. We then filter these boxes by some threshold of area and aspect ratio. Since license plates are typically white in background, and some other color for the text, we further filter these returned contours by the ratio of white pixels to total pixels in that contour. We use a bound of 0.45 as our second to last filter. Finally, for each of these last possible bounding boxes, we pass each to our segmentation algorithm that decomposes each character of a license plate into a cropped image with just that character. If the number of cropped character images returned by the segmentation algorithm is between three and eight inclusive, the majority of the time that is the true license plate. All other bounding boxes are eliminated.

3.3. Character Segmentation

Initially, we approached the task by using a neural network on the entire plates, but this produced bad results. This made sense, because the task was quite complex given the variance of license plates in our dataset. We then tried implementing segmentation of plates into characters to simplify

108 the overall task of getting the plate number to just the task
 109 of identifying one character at a time.
 110

111 To do this, we first cleaned the images to a more easily
 112 processed version. We made the plate image greyscale,
 113 resized and applied a Gaussian blur, and equalized the his-
 114 togram. By then applying a threshold, we got our cleaned im-
 115 age. The main algorithm involved applying a bitwise not on
 116 the cleaned image, finding the contours, and going through
 117 the contours to process each. On each, we found the bound-
 118 ing rectangle, found the area of the bounding box, and de-
 119 pending on whether the area was too small or large, we
 120 dropped the contour as needed with thresholds of 0.008 and
 121 0.032 respectively relative to our entire image area. Other-
 122 wise, we drew a filled rectangle corresponding to the bound-
 123 ing box on the mask with white colour and stored the bound-
 124 ing box/image. The following code highlights our approach
 125 for each contour:

```
126 x,y,w,h = cv2.boundingRect(contour)
127 area = cv2.contourArea(contour)
128 center = (x + w/2, y + h/2)
129 if (float(area/img_area) > 0.008) and /
130   (float(area/img_area) < 0.032) and h > w:
131   x,y,w,h = x-4, y-4, w+8, h+8
132   bounding_boxes.append((center, (x,y,w,h)))
133   cv2.rectangle(char_mask, (x,y), /
134     (x+w,y+h), 255, -1)
```

134 This approach to image character segmentation allowed us
 135 to tackle the task of recognizing each character for each
 136 separated bounding box and image using our neural network
 137 with higher reliability than the initial approach of just one
 138 neural net.

3.4. Neural Network

141 Once the characters of a license plate are segmented into
 142 a cropped 100 by 50 pixel grayscale image, we pass them to
 143 a neural network to predict what digit or letter the character
 144 is. We used Tensorflow to create our model, which takes in a
 145 4D tensor representing a batch of images that are 100 by 50
 146 pixels with a single color channel and returns a probability
 147 distribution over 36 possible characters (10 digits and 26
 148 letters) for each image in the batch. Our model architecture
 149 is as follows:

- 150 • Convolution layer with 32 filters of size 2, a stride of
 151 1, same padding, and ReLU activation
- 152 • Maxpooling layer of 2 by 2
- 153 • Flattening layer
- 154 • Dense layer of size 512 with ReLU activation
- 155 • Dropout layer of 0.5
- 156 • Dense layer of size 36 with softmax activation

162 We train our model for 6 epochs, batch size of 32, a learn-
 163 ing rate of 0.001, Adam optimizer, and SparseCategorical-
 164 Crossentropy loss.

4. Results

4.1. Plate Detector

165 Without any discrete labels of correct bounding boxes,
 166 we were not able to numerically quantify the accuracy of our
 167 detector. However, from user testing, our algorithm seems
 168 fairly reliable. Figure 1 contains one example of the output
 169 before and after filtering. The left image contains the actual
 170 license plate and a returned bounding box but several non-
 171 sensical boxes as well. After passing each of these bounding
 172 boxes to our segmentation algorithm, we only maintain the
 173 bounding boxes that return three to eight segmented char-
 174 acters. In this case, only the license plate holds, so it is the
 175 only bounding box returned. We then crop this bounding
 176 box and pass its segmented characters to our neural network
 177 to read a prediction. However, one issue we ran into is that
 178 the cropped image returned by the detector is much lower
 179 in quality. This may lead to incorrect readings. For instance,
 180 the returned license plate returned in Figure 1 is actually
 181 YG9X2G but is read by our network as Y09X20. The two G
 182 characters are replaced with zeros likely because of image
 183 quality loss as well as the fact that our train dataset contained
 184 many more digits than letters. This necessary filtering and
 185 image quality loss makes the detector is the least reliable out
 186 of the three main steps of our project.

4.2. Character Segmentation

187 Similar to the detector, we did not have any labeled data
 188 of correct bounding boxes for segmentation, so we had to
 189 evaluate the segmentation algorithm by hand for the most
 190 part. This consisted of tests on many different cropped li-
 191 cense plates and saving the segmented results to disc and
 192 examining each segmented image for reliability. An example
 193 of this decomposition can be seen in Figure 2, where we
 194 have the bounding boxes returned for each character for a
 195 single license plate and the first nine of that same license
 196 plate cropped and filtered as it will be passed into our neural
 197 network.

4.3. Neural Network

198 With a train set of size about 3000 and test set of about
 199 300, we were able to get fairly decent results among the 36
 200 possible character classes. We trained for our model for 6
 201 epochs and our plot of loss overtime is shown in Figure 3.
 202 As can be shown, our model converges. For each epoch, we
 203 printed the current model's accuracy on the test set. These re-
 204 sults are shown below. We were able to reach a peak accuracy
 205 of 78.4% on the test set. Epoch results:

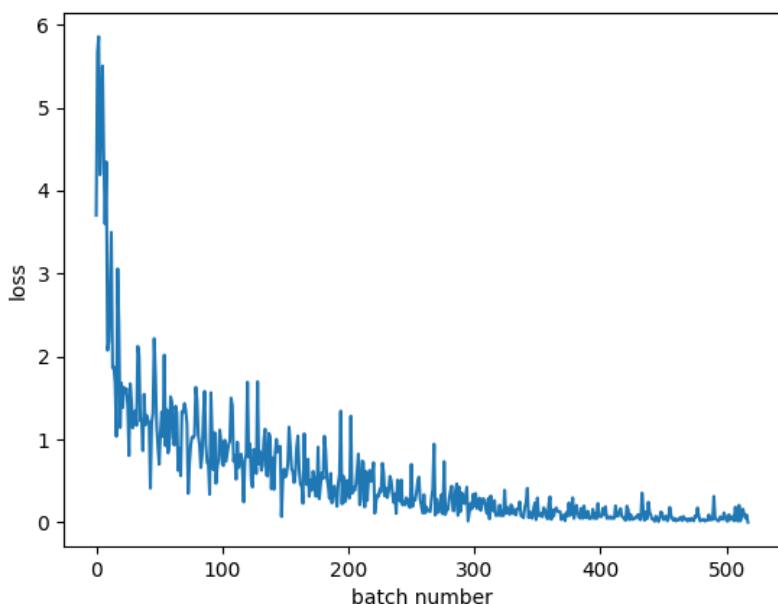
216
217 1 EPOCH 0 | ACCURACY: 0.75769
218 2 EPOCH 1 | ACCURACY: 0.78461
219 3 EPOCH 2 | ACCURACY: 0.71923
220 4 EPOCH 3 | ACCURACY: 0.69999
221 5 EPOCH 4 | ACCURACY: 0.73076
222 6 EPOCH 5 | ACCURACY: 0.76538
223
224



225
226
227
228
229
230
231
232 Figure 1. *Left:* Boxes returned before segmentation filtering. *Center:* After segmentation filtering. *Right:* Lower quality cropped plate
233
234



235
236
237
238
239
240
241
242 Figure 2. *Left:* Bounding boxes for characters. *Right:* Cropped and filtered first character.
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269



267
268
269 Figure 3. Graph of loss overtime
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323

324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377

4.4. Discussion

Some of the components of our method had a variety of approaches that can be discussed. One is the setup of segmenting the license plates into their respective characters, rather than looking at the license plate as a whole, which was our initial approach. One trade-off was the speed and power needed for our program; processing each image in this way used more resources than the first method, and made it take longer to run and more difficult to ensure correctness. However, the benefit seemed to outweigh the cost, since doing this helped our model's accuracy a great deal by simplifying each individual task, despite there being more tasks overall (since there could be up to 8 image recognition values per plate). Our preprocess was also a big part of this improved accuracy, as segmentation allowed for each character to be extracted with more image clarity as opposed to entering an entire license plate image in which there are many variations throughout examples. We can think of this question as a choice between correctness and efficiency, which is a common issue in Deep Learning tasks.

Another part of our method that raised an interesting question with regards to the results was our model. It is difficult to determine a level that real applications of license plate detection would require to be considered reliable. We were able to improve our model through experimentation to high levels for an initial project, with accuracy across epochs upwards of 75%. With additional resources and more extensive datasets, it could be reasonable to aim for much higher accuracy so a model can not only read the characters correctly but the entire license plate. Our model was also designed with a tradeoff between efficiency and correctness. For the purposes of this project, we aimed to have a relatively simple model that could be trained quickly on a laptop. We did not want to overcomplicate it with additional parameters beyond what was needed to give good accuracy on the character model. In

Deep Learning, this decision of simplifying a model versus adding potentially helpful layers is present in almost every task.

5. Conclusion

Thus, our project aimed to process and read license plates in a way that was very applicable to many real-world issues. By using edge detection, image processing, and a deep learning model, we were able to tackle the task of automating human license plate readings with a computer vision program.

This kind of program can prove very valuable moving forward, especially when fine tuned to high accuracy. The project demonstrates potential to add automation to traffic systems and law enforcement by giving humans a reliable way to quickly identify license plates involved in road incidents. This task has the potential to reduce much manual human effort in license plate management by simply requiring an image of the car to get the necessary information.

Appendix

Team contributions

Prithu Dasgupta Contributed to the plate detection, model fine-tuning, and neural network setup.

Daniel Park Contributed to the initial model setup, data gathering and preprocessing, and report.

Gokul Ajith Contributed to the data preprocessing, character segmentation, model training, and command-line processing.

Spencer Greene Contributed to the segmented dataset, character segmentation, and model optimization.

378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431